

Manikantan (Mani) Srinivasan

Boston, MA | +1 (857)-218-2285 | mani.srinivasan2k1@gmail.com | [LinkedIn](#) | [Portfolio](#) | [GitHub](#)

PROFESSIONAL EXPERIENCE

Nebius Jan 2026 – Present
New York City, NY
Forward Deployed Engineer

- Drove enterprise adoption of Tavily, Nebius's AI search product, converting 40+ enterprise clients into paying customers and contributing to \$1M+ in revenue by enabling production integrations and resolving technical adoption barriers.
- Architected and deployed agentic automation pipelines using Claude Agent SDK, LangGraph orchestration, and frontier OpenAI models, integrating with HubSpot CRM and Omni analytics, reducing client response times by 20% and automating internal support workflows.
- Evangelized Tavily's AI search capabilities across enterprise ecosystems, presenting at multiple AI conferences and industry events and hosting technical webinars, while collaborating with partnerships, customer success, and product teams to engage engineers from organizations such as JPMorgan Chase and IBM, generating 500+ enterprise leads.
- Diagnosed and resolved 50+ production issues with core engineering teams using Python, FastAPI, MongoDB, and Snowflake, delivering bug fixes and feature improvements. Improved platform reliability and accelerated enterprise deployments.
- Benchmarked open-source LLMs (e.g., GPT-OSS) by integrating Tavily Search into evaluation pipelines using SEAL-QA and BrowseComp datasets.

SignoFi Sep 2024 – Jan 2025
Raleigh, NC
AI Engineer Intern

- Pioneered an Agentic AI application leveraging Llama 70B and LangChain to retrieve balance sheet data from NoSQL databases, enabling interactive queries and computational analysis for over 50,000 public companies according to regulatory boards.
- Leveraged Fine-tuning with adaptive chunking strategies to ground Llama models, boosting the Faithfulness score to 0.84 (84%) and enhancing contextual precision, recall, and relevancy (92%) in factual question answering on the OpenAI API.
- Architected a containerized custom workflow on Google Cloud Platform using Docker CI/CD pipelines, Vertex AI, FAISS, GCP Buckets, and multi-core GPUs to optimize operations at a portfolio management firm, achieving up to 30% in cost savings.

KEK Aug 2022 – Aug 2023
Remote
Software Engineer – R&D

- Developed production-level distributed computing solutions for the prestigious BELLE-2 particle acceleration experiment, officially part of KEK and working across 5+ world-renowned laboratories including KEK, Brookhaven and CERN.
- Created innovative software tools that empowered physicists to experiment with over 10 PB of data across centers.
- Researched on grid computing techniques and Graph ML to help physicists interpret particle data across 127+ institutions.

EDUCATION

Northeastern University Sep 2023 – Dec 2025
Khoury College of Computer Sciences
MS in Artificial Intelligence - GPA: 3.7
(Expected)
Boston, MA
Relevant Coursework: NLP, AI in Human Computer Interaction, LLM, Machine Learning, Data Science, Unsupervised ML
Teaching Assistant for CS6220 Data Mining – created assignment sets for graduate-level data mining course for topics like ARM, PCA, Data visualization and Analytics.
Research Area– Generative Multimodal AI and Agentic Systems – MCP, Diffusion Transformers, LLMs and Agentic Analytics

SKILLS

Languages: Python, Java, C, C++, Go, HTML, CSS, JavaScript, Swift, Scala, R
Databases & Warehouses: MongoDB, ArangoDB, SQL, NoSQL, Postgres, VectorDB, Redis, Pinecone, GraphQL, Snowflake
Cloud & DevOps and SWE: CI/CD, Docker, AWS, OOPs, AWS S3, GCP, Azure, Git, Agile development, GitHub, GitLab
Data Science: Pandas, Plotly, Seaborn, NumPy, MATLAB, KDP, HubSpot, Omni
AIML: PyTorch, HuggingFace, Tensorflow, LangChain, LangSmith, Claude Agent SDK, A2A, MCP, Reinforcement Learning, Fine-tuning

PROJECTS

ClinAI: MCP-based Agentic platform for voice enabled patient record management

- Engineered an end-to-end agentic AI clinical assistant (ClinAI) leveraging Model Context Protocol and Gemini LLMs to automate extraction and structuring of clinical data from conversations, reducing manual documentation time for doctors by over 60%.
- Deployed a multi-agent workflow pipeline for real-time patient record updating and action item extraction, achieving a low categorisation-error rate of 6.2% on voice transcriptions and reducing manual documentation workload across 100+ clinical transcripts. - <https://github.com/mani2001/clinai-agent>

Machine Unlearning: Erase key concepts from LLM's Knowledge base

- Designed and applied 3 machine unlearning techniques (SAE, LoRA Fine-tuning, Guardrail) on a 7B parameter LLM, ensuring a retained perplexity score under 20 for unaffected knowledge during concept forgetting in HuggingFace Transformers.
- Performed a comparative analysis with both qualitative and perplexity-based studies. Performed hyper-parameter tuning for s with 1B+ parameters. - <https://github.com/mani2001/Machine-Unlearning>